

# AdaBoost

Jiri Matas and Jan Šochman

Centre for Machine Perception  
Czech Technical University, Prague  
<http://cmp.felk.cvut.cz>



## Outline:

- ◆ AdaBoost algorithm
  - Why is of interest?
  - How it works?
  - Why it works?
- ◆ AdaBoost variants
- ◆ AdaBoost with a Totally Corrective Step (TCS)
- ◆ Experiments with a Totally Corrective Step

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

- ◆ 1990 – Boost-by-majority algorithm (Freund)
- ◆ 1995 – AdaBoost (Freund & Schapire)
- ◆ 1997 – Generalized version of AdaBoost (Schapire & Singer)
- ◆ 2001 – AdaBoost in Face Detection (Viola & Jones)

Interesting properties:

- ◆ AB is a **linear** classifier with all its desirable properties.
- ◆ AB output **converges** to the logarithm of likelihood ratio.
- ◆ AB has **good generalization** properties.
- ◆ AB is a **feature selector** with a principled strategy (minimisation of upper bound on empirical error).
- ◆ AB **close to sequential decision making** (it produces a sequence of gradually more complex classifiers).

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

- AdaBoost is an algorithm for constructing a "strong" classifier as linear combination

$$f(x) = \sum_{t=1}^T \alpha_t h_t(x)$$

of "simple" "weak" classifiers  $h_t(x)$ .

## Terminology

- $h_t(x)$  ... "weak" or basis classifier, hypothesis, "feature"
- $H(x) = \text{sign}(f(x))$  ... "strong" or final classifier/hypothesis

## Comments

- The  $h_t(x)$ 's can be thought of as features.
- Often (typically) the set  $\mathcal{H} = \{h(x)\}$  is infinite.

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

Given:  $(x_1, y_1), \dots, (x_m, y_m); x_i \in \mathcal{X}, y_i \in \{-1, 1\}$

Initialize weights  $D_1(i) = 1/m$

For  $t = 1, \dots, T$ :

1. (Call *WeakLearn*), which returns the weak classifier  $h_t : \mathcal{X} \rightarrow \{-1, 1\}$  with minimum error w.r.t. distribution  $D_t$ ;
2. Choose  $\alpha_t \in R$ ,
3. Update

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

where  $Z_t$  is a normalization factor chosen so that  $D_{t+1}$  is a distribution

Output the strong classifier:

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$$

### Comments

- ◆ The computational complexity of selecting  $h_t$  is independent of  $t$ .
- ◆ All information about previously selected “features” is captured in  $D_t$ !

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

**Loop step:** Call *WeakLearn*, given distribution  $D_t$ ;  
returns weak classifier  $h_t : \mathcal{X} \rightarrow \{-1, 1\}$  from  $\mathcal{H} = \{h(x)\}$

- ◆ Select a weak classifier with the smallest weighted error

$$h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i)[y_i \neq h_j(x_i)]$$

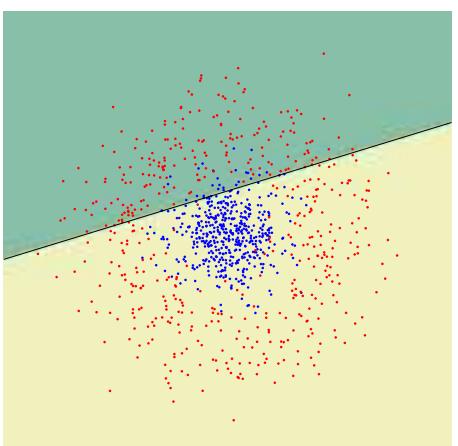
- ◆ Prerequisite:  $\epsilon_t < 1/2$  (otherwise stop)

- ◆ *WeakLearn* examples:

- Decision tree builder, perceptron learning rule –  $\mathcal{H}$  infinite
- Selecting the best one from given *finite* set  $\mathcal{H}$

## Demonstration example

Training set



Weak classifier = perceptron

- $\sim N(0, 1)$
- $\sim \frac{1}{r\sqrt{8\pi^3}} e^{-1/2(r-4)^2}$

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16



- ◆ The main objective is to minimize  $\varepsilon_{tr} = \frac{1}{m} |\{i : H(x_i) \neq y_i\}|$
- ◆ It can be upper bounded by  $\varepsilon_{tr}(H) \leq \prod_{t=1}^T Z_t$

## How to set $\alpha_t$ ?

- ◆ Select  $\alpha_t$  to greedily minimize  $Z_t(\alpha)$  in each step
- ◆  $Z_t(\alpha)$  is convex differentiable function with one extremum  
 $\Rightarrow h_t(x) \in \{-1, 1\}$  then optimal  $\alpha_t = \frac{1}{2} \log(\frac{1+r_t}{1-r_t})$   
where  $r_t = \sum_{i=1}^m D_t(i) h_t(x_i) y_i$
- ◆  $Z_t = 2\sqrt{\epsilon_t(1-\epsilon_t)} \leq 1$  for optimal  $\alpha_t$   
 $\Rightarrow$  Justification of selection of  $h_t$  according to  $\epsilon_t$

## Comments

- ◆ The process of selecting  $\alpha_t$  and  $h_t(x)$  can be interpreted as a single optimization step minimising the upper bound on the empirical error.  
Improvement of the bound is guaranteed, provided that  $\epsilon_t < 1/2$ .
- ◆ The process can be interpreted as a component-wise local optimization (Gauss-Southwell iteration) in the (possibly infinite dimensional!) space of  $\bar{\alpha} = (\alpha_1, \alpha_2, \dots)$  starting from  $\bar{\alpha}_0 = (0, 0, \dots)$ .

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

## Effect on the training set

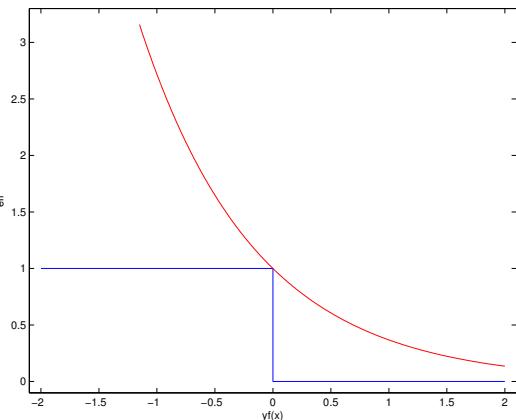
Reweighting formula:

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t} = \frac{\exp(-y_i \sum_{q=1}^t \alpha_q h_q(x_i))}{m \prod_{q=1}^t Z_q}$$

$$\exp(-\alpha_t y_i h_t(x_i)) \begin{cases} < 1, & y_i = h_t(x_i) \\ > 1, & y_i \neq h_t(x_i) \end{cases}$$

}

⇒ Increase (decrease) weight of wrongly (correctly) classified examples. The weight is the upper bound on the error of a given example!



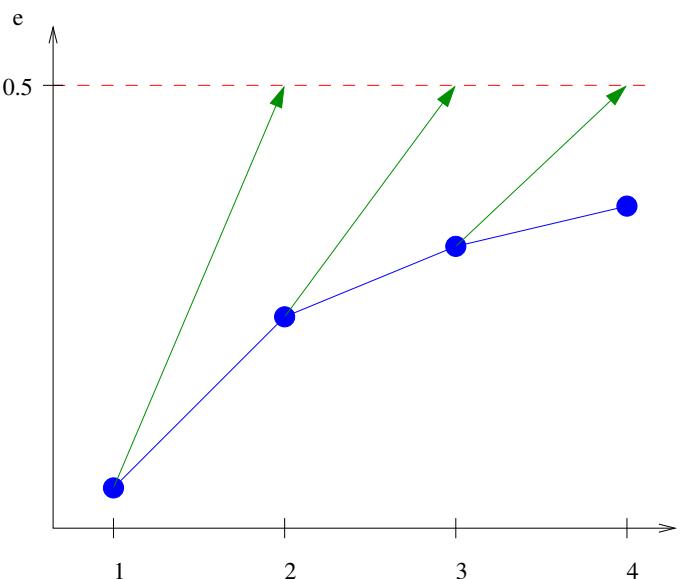
## Effect on $h_t$

◆  $\alpha_t$  minimize  $Z_t \Rightarrow$

$$\sum_{i:h_t(x_i)=y_i} D_{t+1}(i) = \sum_{i:h_t(x_i) \neq y_i} D_{t+1}(i)$$

◆ Error of  $h_t$  on  $D_{t+1}$  is  $1/2$

◆ Next weak classifier is the most “independent” one



1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

# Summary of the Algorithm



m p

Initialization...

For  $t = 1, \dots, T$ :

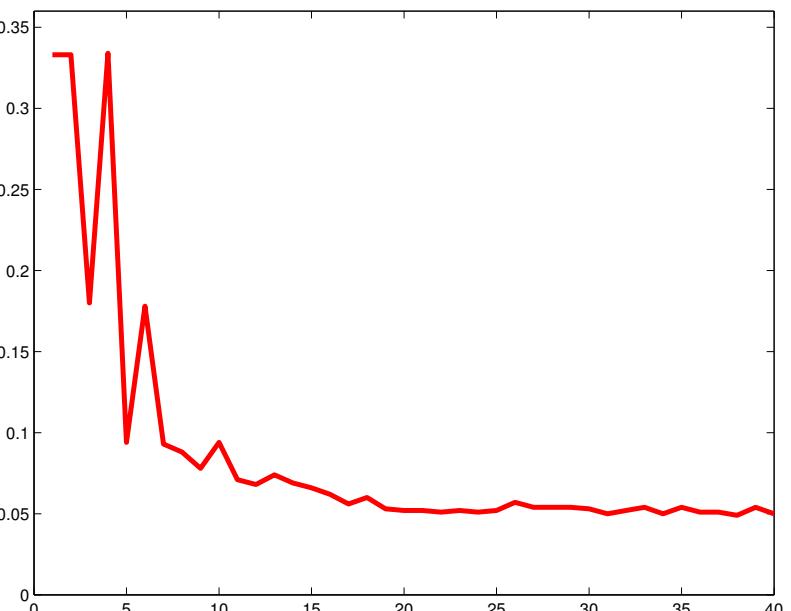
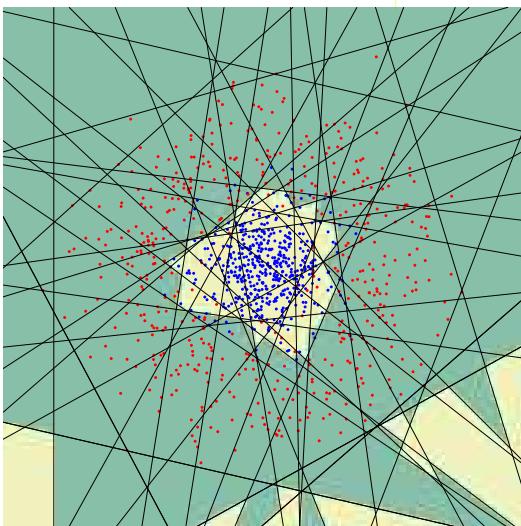
- ◆ Find  $h_t = \arg \min_{h_j \in \mathcal{H}} \epsilon_j = \sum_{i=1}^m D_t(i)[y_i \neq h_j(x_i)]$
- ◆ If  $\epsilon_t \geq 1/2$  then stop
- ◆ Set  $\alpha_t = \frac{1}{2} \log(\frac{1+r_t}{1-r_t})$
- ◆ Update

$$D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

Output the final classifier:

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$$

$t = 40$



1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

Margins in SVM

$$\max \min_{(x,y) \in S} \frac{y(\vec{\alpha} \cdot \vec{h}(x))}{\|\vec{\alpha}\|_2}$$

Margins in AdaBoost

$$\max \min_{(x,y) \in S} \frac{y(\vec{\alpha} \cdot \vec{h}(x))}{\|\vec{\alpha}\|_1}$$

**Maximizing margins in AdaBoost**

$$P_S[yf(x) \leq \theta] \leq 2^T \prod_{t=1}^T \sqrt{\epsilon_t^{1-\theta} (1-\epsilon_t)^{1+\theta}} \quad \text{where } f(x) = \frac{\vec{\alpha} \cdot \vec{h}(x)}{\|\vec{\alpha}\|_1}$$

**Upper bounds based on margin**

$$P_{\mathcal{D}}[yf(x) \leq 0] \leq P_S[yf(x) \leq \theta] + \mathcal{O} \left( \frac{1}{\sqrt{m}} \left( \frac{d \log^2(m/d)}{\theta^2} + \log(1/\delta) \right)^{1/2} \right)$$

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

## Freund & Schapire 1995

- ◆ Discrete ( $h : \mathcal{X} \rightarrow \{0, 1\}$ )
- ◆ Multiclass AdaBoost.M1 ( $h : \mathcal{X} \rightarrow \{0, 1, \dots, k\}$ )
- ◆ Multiclass AdaBoost.M2 ( $h : \mathcal{X} \rightarrow [0, 1]^k$ )
- ◆ Real valued AdaBoost.R ( $Y = [0, 1]$ ,  $h : \mathcal{X} \rightarrow [0, 1]$ )

## Schapire & Singer 1997

- ◆ Confidence rated prediction ( $h : \mathcal{X} \rightarrow R$ , two-class)
- ◆ Multilabel AdaBoost.MR, AdaBoost.MH (different formulation of minimized loss)

... Many other modifications since then (Totally Corrective AB, Cascaded AB)

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16



## Advantages

- ◆ Very simple to implement
- ◆ Feature selection on very large sets of features
- ◆ Fairly good generalization

## Disadvantages

- ◆ Suboptimal solution for  $\bar{\alpha}$
- ◆ Can overfit in presence of noise

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

Given:  $(x_1, y_1), \dots, (x_m, y_m); x_i \in \mathcal{X}, y_i \in \{-1, 1\}$

Initialize weights  $D_1(i) = 1/m$

For  $t = 1, \dots, T$ :

1. (Call *WeakLearn*), which returns the weak classifier  $h_t : \mathcal{X} \rightarrow \{-1, 1\}$  with minimum error w.r.t. distribution  $D_t$ ;
2. Choose  $\alpha_t \in R$ ,
3. Update  $D_{t+1}$
4. (Call *WeakLearn*) on the set of  $h_m$ 's with non zero  $\alpha$ 's . Update  $\alpha$ ..  
Update  $D_{t+1}$ . Repeat till  $|\epsilon_t - 1/2| < \delta, \forall t$ .

## Comments

- ◆ All weak classifiers have  $\epsilon_t \approx 1/2$ , therefore the classifier selected at  $t + 1$  is "independent" of all classifiers selected so far.
- ◆ It can be easily shown, that the totally corrective step reduces the upper bound on the empirical error without increasing classifier complexity.
- ◆ The TCA was first proposed by Kivinen and Warmuth, but their  $\alpha_t$  is set as in standard Adaboost.
- ◆ Generalization of TCA is an open question.

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

- ◆ Discrete AdaBoost, Real AdaBoost, and Discrete and Real TCA evaluated
- ◆ Weak learner: stumps.
- ◆ Data from the IDA repository (Ratsch:2000):

	Input dimension	Training patterns	Testing patterns	Number of realizations
Banana	2	400	4900	100
Breast cancer	9	200	77	100
Diabetes	8	468	300	100
German	20	700	300	100
Heart	13	170	100	100
Image segment	18	1300	1010	20
Ringnorm	20	400	7000	100
Flare solar	9	666	400	100
Splice	60	1000	2175	20
Thyroid	5	140	75	100
Titanic	3	150	2051	100
Twonorm	20	400	7000	100
Waveform	21	400	4600	100

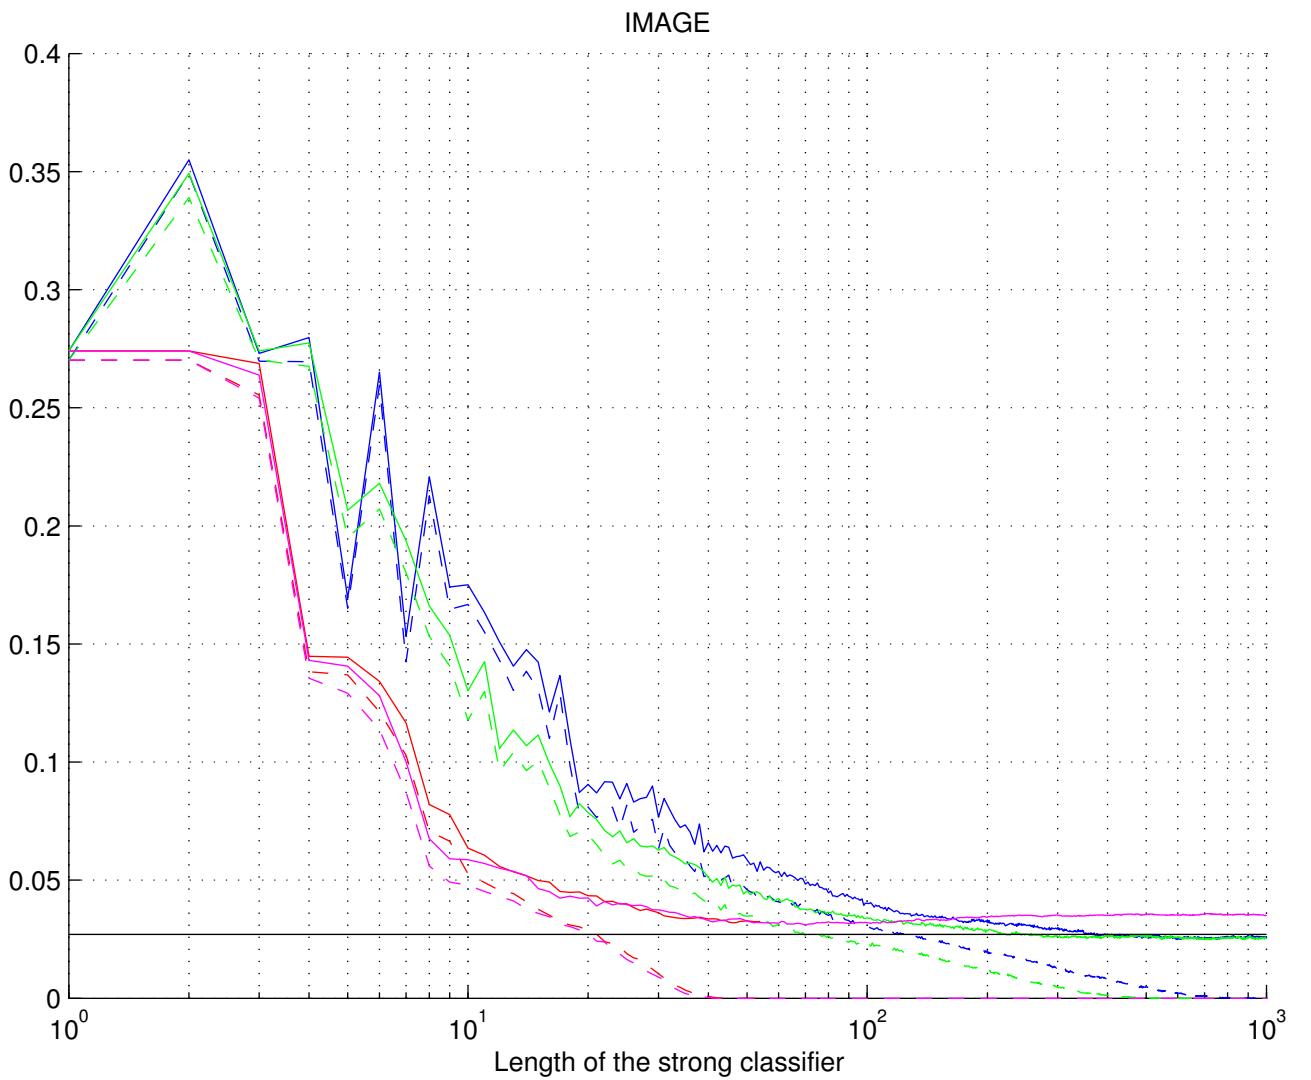
- ◆ Note that the training sets are fairly small

# Results with TCA on the IDA Database



m p

- ◆ Training error (dashed line), test error (solid line)
- ◆ Discrete AdaBoost (blue), Real AdaBoost (green),
- ◆ Discrete AdaBoost with TCA (red), Real AdaBoost with TCA (cyan)
- ◆ the black horizontal line: the error of AdaBoost with RBF network weak classifiers from (Ratsch-ML:2000)



16

- ◆ The AdaBoost algorithm was presented and analysed
- ◆ A modification of the Totally Corrective AdaBoost was introduced
- ◆ Initial test show that the TCA outperforms AB on some standard data sets.

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16